

Denoising Diffusion Probabilistic Models (DDPMs)

**Presented by
MEDIocre_GUY**

April 4, 2024




Introduction (1/2)

In simple terms, two processes happen in denoising diffusion probabilistic models (DDPMs):

- The data structure is destroyed by gradually adding Gaussian noise over a finite number of time steps to end up with pure noise (forward/diffusion process)
- A neural network is trained to gradually denoise the data starting from pure noise and predict a distribution that looks like the original distribution (reverse/denoising process)

Forward/diffusion process




Reverse/denoising process





Introduction (2/2)

The objective of using diffusion models is to successfully generate actual images from pure noise only if they are trained well



Source: https://github.com/TeaPearce/Conditional_Diffusion_MNIST

Forward (Diffusion) Process (1/2)

$q(x_0)$ = Original data distribution

$x_0 \sim q(x_0)$



Taking a sample from the original data distribution

$$q(x_t|x_{t-1}) = \mathcal{N}(x_t; \sqrt{1 - \beta_t}x_{t-1}, \beta_t\mathbf{I})$$

Forward process

Mean (μ_t): $\sqrt{1 - \beta_t}x_{t-1}$

Variance (σ_t^2): β_t

x_{t-1} = Less noisy image

x_t = More noisy image

β_t = *Variance scheduler* (linear, cosine, sigmoid, quadratic etc.)

Forward process $q(x_t|x_{t-1})$ adds Gaussian noise $\varepsilon \sim \mathcal{N}(\mathbf{0}, \mathbf{I})$ according to a known variance schedule ($0 < \beta_t < 1$)

β_t are constants that increase over T time steps

Original DDPM paper used *linear* scheduler ($\beta_1 = 0.0001$ to $\beta_T = 0.02$ for $T = 1000$)

The source image (x_0) eventually turns into pure noise (x_T) *through the forward process*



Forward (Diffusion) Process (2/2)

$q(x_t|x_{t-1}) = \mathcal{N}(x_t; \sqrt{1 - \beta_t}x_{t-1}, \beta_t\mathbf{I}) \longrightarrow$ A single step of the forward process

$q(x_{1:T}|x_0) = \prod_{t=1}^T q(x_t|x_{t-1}) \longrightarrow$ Equation for the *full forward process*

Reparameterization trick $\longrightarrow \mathcal{N}(\mu, \sigma^2) = \mu + \sigma \odot \varepsilon$

$\odot \rightarrow$ Element-wise product

$$x_t = \sqrt{1 - \beta_t}x_{t-1} + \sqrt{\beta_t}\varepsilon \quad \varepsilon \sim \mathcal{N}(\mathbf{0}, \mathbf{I})$$

$$= \sqrt{\alpha_t}x_{t-1} + \sqrt{1 - \alpha_t}\varepsilon \quad \alpha_t = 1 - \beta_t$$

$$= \sqrt{\alpha_t\alpha_{t-1}}x_{t-2} + \sqrt{1 - \alpha_t\alpha_{t-1}}\varepsilon$$

.....

$$= \sqrt{\bar{\alpha}_t}x_0 + \sqrt{1 - \bar{\alpha}_t}\varepsilon$$

$$\bar{\alpha}_t = \prod_{s=1}^t \alpha_s$$

For example, $\bar{\alpha}_3 = \alpha_1 \cdot \alpha_2 \cdot \alpha_3$

$$x_t = \sqrt{\bar{\alpha}_t}x_0 + \sqrt{1 - \bar{\alpha}_t}\varepsilon$$

$$q(x_t|x_0) = \mathcal{N}(x_t; \sqrt{\bar{\alpha}_t}x_0, (1 - \bar{\alpha}_t)\mathbf{I})$$

This allows to sample x_t at any time step t conditioned on x_0



Reverse (Denoising) Process (1/15)

$$p(x_{t-1}|x_t)$$

**Reverse
process**



$$p_{\theta}(x_{t-1}|x_t) \quad \theta \rightarrow \text{all the parameters of the neural network}$$

$$p_{\theta}(x_{t-1}|x_t) = \mathcal{N}(x_{t-1}; \mu_{\theta}(x_t, t), \Sigma_{\theta}(x_t, t))$$

Both the mean (μ_{θ}) and the variance (Σ_{θ}) are conditioned on the noise level (time step) t

In the original DDPM paper, the authors kept the variance (Σ_{θ}) fixed and used one neural network to learn only the mean (μ_{θ})

$$\Sigma_{\theta}(x_t, t) = \sigma_t^2 \mathbf{I}$$

$$\sigma_t^2 = \beta_t$$

$$\sigma_t^2 = \tilde{\beta}_t = \frac{1 - \bar{\alpha}_{t-1}}{1 - \bar{\alpha}_t} \beta_t$$

DDPM authors claimed that both choices yielded *similar results*



Reverse (Denoising) Process (2/15)

To learn the mean (μ_θ), the DDPM authors decided to minimize the negative log-likelihood (NLL) regarding the source image as their *objective function*

↳ $-\log(p_\theta(x_0))$

↳ **Not computable**

Have to keep track of $T - 1$ other random variables
($x_T, x_{T-1}, x_{T-2}, \dots, x_3, x_2, x_1$)

So, the DDPM authors chose to compute the *variational lower bound (VLB)* instead

↓

$$-\log(p_\theta(x_0)) \leq -\log(p_\theta(x_0)) + D_{KL}(q(x_{1:T}|x_0) || p_\theta(x_{1:T}|x_0))$$

KL divergence:

$$D_{KL}(p||q) = \sum_x p(x) \log \left(\frac{p(x)}{q(x)} \right)$$

Needs to be minimized



Reverse (Denoising) Process (3/15)

$$D_{KL}(q(x_{1:T}|x_0)||p_{\theta}(x_{1:T}|x_0)) \longrightarrow \log \left(\frac{q(x_{1:T}|x_0)}{p_{\theta}(x_{1:T}|x_0)} \right)$$

Joint probability

$$p_{\theta}(x_{1:T}|x_0) = \underbrace{\frac{p_{\theta}(x_0|x_{1:T})p_{\theta}(x_{1:T})}{p_{\theta}(x_0)}}_{\text{Bayes' rule}} = \frac{\overbrace{p_{\theta}(x_0, x_{1:T})}^{\text{Joint probability}}}{p_{\theta}(x_0)} = \frac{p(x_{0:T})}{p_{\theta}(x_0)}$$

Bayes' rule

$$\log \left(\frac{q(x_{1:T}|x_0)}{p_{\theta}(x_{1:T}|x_0)} \right) = \log \left(\frac{q(x_{1:T}|x_0)}{\frac{p(x_{0:T})}{p_{\theta}(x_0)}} \right) = \underbrace{\log \left(\frac{q(x_{1:T}|x_0)}{p(x_{0:T})} \right) + \log(p_{\theta}(x_0))}_{\text{Log product rule}}$$

Log product rule

$$-\log(p_{\theta}(x_0)) \leq -\log(p_{\theta}(x_0)) + D_{KL}(q(x_{1:T}|x_0)||p_{\theta}(x_{1:T}|x_0)) \quad \text{now becomes}$$



Reverse (Denoising) Process (4/15)

$$-\log(p_\theta(x_0)) \leq \boxed{-\log(p_\theta(x_0))} + \log\left(\frac{q(x_{1:T}|x_0)}{p_\theta(x_{0:T})}\right) \boxed{+\log(p_\theta(x_0))}$$

$$-\log(p_\theta(x_0)) \leq \log\left(\frac{q(x_{1:T}|x_0)}{p_\theta(x_{0:T})}\right) \quad \text{Variational Lower Bound (VLB)}$$

$q(x_{1:T}|x_0)$ is the *forward process*

$$p_\theta(x_{0:T}) = p(x_T) \prod_{t=1}^T p_\theta(x_{t-1}|x_t)$$

$$\log\left(\frac{q(x_{1:T}|x_0)}{p_\theta(x_{0:T})}\right) = \log\left(\frac{\prod_{t=1}^T q(x_t|x_{t-1})}{p(x_T) \prod_{t=1}^T p_\theta(x_{t-1}|x_t)}\right) = -\log(p(x_T)) + \sum_{t=1}^T \log\left(\frac{q(x_t|x_{t-1})}{p_\theta(x_{t-1}|x_t)}\right)$$


$$= -\log(p(x_T)) + \sum_{t=2}^T \log\left(\frac{q(x_t|x_{t-1})}{p_\theta(x_{t-1}|x_t)}\right) + \log\left(\frac{q(x_1|x_0)}{p_\theta(x_0|x_1)}\right)$$



Reverse (Denoising) Process (5/15)

$$q(x_t|x_{t-1}) = \frac{q(x_{t-1}|x_t)q(x_t)}{q(x_{t-1})} \quad \text{According to Bayes' rule}$$

The DDPM authors
decided to condition this
on x_0 (trick)



$$\frac{q(x_{t-1}|x_t, x_0)q(x_t|x_0)}{q(x_{t-1}|x_0)}$$

So, VLB becomes

$$\begin{aligned} & -\log(p(x_T)) + \sum_{t=2}^T \log\left(\frac{q(x_{t-1}|x_t, x_0)q(x_t|x_0)}{p_\theta(x_{t-1}|x_t)q(x_{t-1}|x_0)}\right) + \log\left(\frac{q(x_1|x_0)}{p_\theta(x_0|x_1)}\right) \\ & = -\log(p(x_T)) + \sum_{t=2}^T \log\left(\frac{q(x_{t-1}|x_t, x_0)}{p_\theta(x_{t-1}|x_t)}\right) + \boxed{\sum_{t=2}^T \log\left(\frac{q(x_t|x_0)}{q(x_{t-1}|x_0)}\right)} + \log\left(\frac{q(x_1|x_0)}{p_\theta(x_0|x_1)}\right) \end{aligned}$$

For, $T = 4$

$$\sum_{t=2}^4 \log\left(\frac{q(x_t|x_0)}{q(x_{t-1}|x_0)}\right) = \log\left(\prod_{t=2}^4 \frac{q(x_t|x_0)}{q(x_{t-1}|x_0)}\right) = \log\left(\frac{\boxed{q(x_2|x_0)q(x_3|x_0)}q(x_4|x_0)}{q(x_1|x_0)\boxed{q(x_2|x_0)q(x_3|x_0)}}\right)$$



Reverse (Denoising) Process (6/15)

Eventually, we obtain

$$\sum_{t=2}^T \log \left(\frac{q(x_t|x_0)}{q(x_{t-1}|x_0)} \right) = \log \left(\frac{q(x_T|x_0)}{q(x_1|x_0)} \right)$$

Hence, VLB becomes

$$-\log(p(x_T)) + \sum_{t=2}^T \log \left(\frac{q(x_{t-1}|x_t, x_0)}{p_\theta(x_{t-1}|x_t)} \right) + \log \left(\frac{q(x_T|x_0)}{q(x_1|x_0)} \right) + \log \left(\frac{q(x_1|x_0)}{p_\theta(x_0|x_1)} \right)$$

Log division rule gives us

$$\log(q(x_T|x_0)) - \log(q(x_1|x_0)) + \log(q(x_1|x_0)) - \log(p_\theta(x_0|x_1))$$

$$-\log(p(x_T)) + \sum_{t=2}^T \log \left(\frac{q(x_{t-1}|x_t, x_0)}{p_\theta(x_{t-1}|x_t)} \right) + \log(q(x_T|x_0)) - \log(p_\theta(x_0|x_1))$$



Reverse (Denoising) Process (7/15)

Simplification results in

$$\log \left(\frac{q(x_T|x_0)}{p(x_T)} \right) + \sum_{t=2}^T \log \left(\frac{q(x_{t-1}|x_t, x_0)}{p_\theta(x_{t-1}|x_t)} \right) - \log(p_\theta(x_0|x_1))$$



$$D_{KL}(q(x_T|x_0)||p(x_T)) + \sum_{t=2}^T D_{KL}(q(x_{t-1}|x_t, x_0)||p_\theta(x_{t-1}|x_t)) - \log(p_\theta(x_0|x_1))$$

This part has no learnable parameters

$$\mathcal{N}(x_{t-1}; \tilde{\mu}_t(x_t, x_0), \tilde{\beta}_t \mathbf{I})$$

$$\mathcal{N}(x_{t-1}; \mu_\theta(x_t, t), \beta_t \mathbf{I})$$

$$q(x_{t-1}|x_t, x_0) = \frac{q(x_t|x_{t-1}, x_0)q(x_{t-1}|x_0)}{q(x_t|x_0)} \quad \text{Bayes' rule}$$



Reverse (Denoising) Process (8/15)

$$\begin{aligned}
 &\propto \exp \left(-\frac{1}{2} \left(\frac{(x_t - \sqrt{\alpha_t} x_{t-1})^2}{\beta_t} + \frac{(x_{t-1} - \sqrt{\bar{\alpha}_{t-1}} x_0)^2}{1 - \bar{\alpha}_{t-1}} - \frac{(x_t - \sqrt{\bar{\alpha}_t} x_0)^2}{1 - \bar{\alpha}_t} \right) \right) \\
 &= \exp \left(-\frac{1}{2} \left(\frac{x_t^2 - 2\sqrt{\alpha_t} x_t x_{t-1} + \alpha_t x_{t-1}^2}{\beta_t} + \frac{x_{t-1}^2 - 2\sqrt{\bar{\alpha}_{t-1}} x_0 x_{t-1} + \bar{\alpha}_{t-1} x_0^2}{1 - \bar{\alpha}_{t-1}} - \frac{(x_t - \sqrt{\bar{\alpha}_t} x_0)^2}{1 - \bar{\alpha}_t} \right) \right) \\
 &= \exp \left(-\frac{1}{2} \left(\left(\frac{\alpha_t}{\beta_t} + \frac{1}{1 - \bar{\alpha}_{t-1}} \right) x_{t-1}^2 - \left(\frac{2\sqrt{\alpha_t}}{\beta_t} x_t + \frac{2\sqrt{\bar{\alpha}_{t-1}}}{1 - \bar{\alpha}_{t-1}} x_0 \right) x_{t-1} + C(x_t, x_0) \right) \right)
 \end{aligned}$$

$$\tilde{\beta}_t = \frac{1}{\frac{\alpha_t}{\beta_t} + \frac{1}{1 - \bar{\alpha}_{t-1}}} = \frac{\beta_t(1 - \bar{\alpha}_{t-1})}{\alpha_t(1 - \bar{\alpha}_{t-1}) + \beta_t} = \frac{1 - \bar{\alpha}_{t-1}}{1 - \bar{\alpha}_t} \beta_t$$

$$\bar{\alpha}_{t-1} = \alpha_1 \cdot \alpha_2 \cdot \alpha_3 \dots \alpha_{t-1}$$

$$\bar{\alpha}_t = \alpha_1 \cdot \alpha_2 \cdot \alpha_3 \dots \alpha_{t-1} \alpha_t$$

$$\beta_t = 1 - \alpha_t$$

$$\bar{\alpha}_t = \prod_{s=1}^t \alpha_s$$

$$\tilde{\beta}_t = \frac{1 - \bar{\alpha}_{t-1}}{1 - \bar{\alpha}_t} \beta_t$$



Reverse (Denoising) Process (9/15)

$$\exp\left(-\frac{1}{2}\left(\left(\frac{\alpha_t}{\beta_t} + \frac{1}{1 - \bar{\alpha}_{t-1}}\right)x_{t-1}^2 - \left(\frac{2\sqrt{\alpha_t}}{\beta_t}x_t + \frac{2\sqrt{\bar{\alpha}_{t-1}}}{1 - \bar{\alpha}_{t-1}}x_0\right)x_{t-1} + C(x_t, x_0)\right)\right)$$
$$\tilde{\mu}_t(x_t, x_0) = \frac{\frac{\sqrt{\alpha_t}}{\beta_t}x_t + \frac{\sqrt{\bar{\alpha}_{t-1}}}{1 - \bar{\alpha}_{t-1}}x_0}{\frac{\alpha_t}{\beta_t} + \frac{1}{1 - \bar{\alpha}_{t-1}}} = \frac{\sqrt{\alpha_t}}{\beta_t}x_t + \frac{\sqrt{\bar{\alpha}_{t-1}}}{1 - \bar{\alpha}_{t-1}}x_0 \left(\frac{1}{\frac{\alpha_t}{\beta_t} + \frac{1}{1 - \bar{\alpha}_{t-1}}}\right)$$
$$= \frac{\sqrt{\alpha_t}}{\beta_t}x_t + \frac{\sqrt{\bar{\alpha}_{t-1}}}{1 - \bar{\alpha}_{t-1}}x_0 \left(\frac{1 - \bar{\alpha}_{t-1}}{1 - \bar{\alpha}_t}\beta_t\right)$$
$$= \frac{\sqrt{\alpha_t}(1 - \bar{\alpha}_{t-1})}{1 - \bar{\alpha}_t}x_t + \frac{\sqrt{\bar{\alpha}_{t-1}}\beta_t}{1 - \bar{\alpha}_t}x_0$$



Reverse (Denoising) Process (10/15)

$$\tilde{\mu}_t(x_t, x_0) = \frac{\sqrt{\alpha_t}(1 - \bar{\alpha}_{t-1})}{1 - \bar{\alpha}_t} x_t + \frac{\sqrt{\bar{\alpha}_{t-1}}\beta_t}{1 - \bar{\alpha}_t} x_0$$

From forward process $\rightarrow x_t = \sqrt{\bar{\alpha}_t}x_0 + \sqrt{1 - \bar{\alpha}_t}\varepsilon$

$$\Rightarrow x_0 = \frac{1}{\sqrt{\bar{\alpha}_t}}(x_t - \sqrt{1 - \bar{\alpha}_t}\varepsilon)$$

$$\begin{aligned}\tilde{\mu}_t(x_t, x_0) &= \frac{\sqrt{\alpha_t}(1 - \bar{\alpha}_{t-1})}{1 - \bar{\alpha}_t} x_t + \frac{\sqrt{\bar{\alpha}_{t-1}}\beta_t}{1 - \bar{\alpha}_t} \frac{1}{\sqrt{\bar{\alpha}_t}}(x_t - \sqrt{1 - \bar{\alpha}_t}\varepsilon) \\ &= \left[\frac{\sqrt{\alpha_t}(1 - \bar{\alpha}_{t-1})}{1 - \bar{\alpha}_t} + \frac{\sqrt{\bar{\alpha}_{t-1}}\beta_t}{\sqrt{\bar{\alpha}_t}(1 - \bar{\alpha}_t)} \right] x_t - \frac{\sqrt{\bar{\alpha}_{t-1}}\beta_t\sqrt{1 - \bar{\alpha}_t}}{\sqrt{\bar{\alpha}_t}(1 - \bar{\alpha}_t)} \varepsilon\end{aligned}$$

$$= Ax_t - B\varepsilon$$



Reverse (Denoising) Process (11/15)

$$\begin{aligned} A &= \frac{\sqrt{\alpha_t}(1 - \bar{\alpha}_{t-1})}{1 - \bar{\alpha}_t} + \frac{\sqrt{\bar{\alpha}_{t-1}}\beta_t}{\sqrt{\bar{\alpha}_t}(1 - \bar{\alpha}_t)} = \frac{\sqrt{\bar{\alpha}_t}\sqrt{\alpha_t}(1 - \bar{\alpha}_{t-1}) + \sqrt{\bar{\alpha}_{t-1}}(1 - \alpha_t)}{\sqrt{\bar{\alpha}_t}(1 - \bar{\alpha}_t)} \quad \beta_t = 1 - \alpha_t \\ &= \frac{\sqrt{\alpha_t \bar{\alpha}_{t-1}}\sqrt{\alpha_t}(1 - \bar{\alpha}_{t-1}) + \sqrt{\bar{\alpha}_{t-1}}(1 - \alpha_t)}{\sqrt{\bar{\alpha}_t}(1 - \bar{\alpha}_t)} \quad \bar{\alpha}_t = \alpha_t \bar{\alpha}_{t-1} \\ &= \frac{\alpha_t \sqrt{\bar{\alpha}_{t-1}} - \alpha_t \sqrt{\bar{\alpha}_{t-1}} \bar{\alpha}_{t-1} + \sqrt{\bar{\alpha}_{t-1}} - \alpha_t \sqrt{\bar{\alpha}_{t-1}}}{\sqrt{\bar{\alpha}_t}(1 - \bar{\alpha}_t)} \\ &= \frac{\sqrt{\bar{\alpha}_{t-1}}(1 - \bar{\alpha}_t)}{\sqrt{\bar{\alpha}_t}(1 - \bar{\alpha}_t)} \quad \bar{\alpha}_t = \alpha_t \bar{\alpha}_{t-1} \\ &= \frac{\sqrt{\bar{\alpha}_{t-1}}}{\sqrt{\alpha_t} \sqrt{\bar{\alpha}_{t-1}}} \quad \bar{\alpha}_t = \alpha_t \bar{\alpha}_{t-1} \\ &= \frac{1}{\sqrt{\alpha_t}} \end{aligned}$$



Reverse (Denoising) Process (12/15)

$$\begin{aligned}
 B &= \frac{\sqrt{\bar{\alpha}_{t-1}}\beta_t\sqrt{1-\bar{\alpha}_t}}{\sqrt{\bar{\alpha}_t}(1-\bar{\alpha}_t)} \\
 &= \frac{\sqrt{\bar{\alpha}_{t-1}}\beta_t\sqrt{1-\bar{\alpha}_t}}{\sqrt{\bar{\alpha}_t}(\sqrt{1-\bar{\alpha}_t})^2} \\
 &= \frac{\sqrt{\bar{\alpha}_{t-1}}\beta_t}{\sqrt{\bar{\alpha}_t}\sqrt{1-\bar{\alpha}_t}} \\
 &= \frac{\sqrt{\bar{\alpha}_{t-1}}\beta_t}{\sqrt{\alpha_t\bar{\alpha}_{t-1}}\sqrt{1-\bar{\alpha}_t}} \\
 &= \frac{\beta_t}{\sqrt{\alpha_t}\sqrt{1-\bar{\alpha}_t}}
 \end{aligned}$$

$$\bar{\alpha}_t = \alpha_t \bar{\alpha}_{t-1}$$

$$\tilde{\mu}_t(x_t, x_0) = Ax_t - B\varepsilon$$

$$A = \frac{1}{\sqrt{\alpha_t}} \quad B = \frac{\beta_t}{\sqrt{\alpha_t}\sqrt{1-\bar{\alpha}_t}}$$

$$\begin{aligned}
 \tilde{\mu}_t(x_t, x_0) &= \frac{1}{\sqrt{\alpha_t}}x_t - \frac{\beta_t}{\sqrt{\alpha_t}\sqrt{1-\bar{\alpha}_t}}\varepsilon \\
 &= \frac{1}{\sqrt{\alpha_t}}\left(x_t - \frac{\beta_t}{\sqrt{1-\bar{\alpha}_t}}\varepsilon\right)
 \end{aligned}$$



Mean of the forward process posterior

$\mu_\theta(x_t, t)$



Mean predicted by the neural network



Reverse (Denoising) Process (13/15)

The distance between $\tilde{\mu}_t(x_t, x_0)$ and $\mu_\theta(x_t, t)$ is approximated using a *mean-squared error (MSE)*

$$L_t = \frac{1}{2\sigma_t^2} \|\tilde{\mu}_t(x_t, x_0) - \mu_\theta(x_t, t)\|^2$$
$$= \frac{1}{2\sigma_t^2} \left\| \frac{1}{\sqrt{\alpha_t}} \left(x_t - \frac{\beta_t}{\sqrt{1 - \bar{\alpha}_t}} \varepsilon \right) - \mu_\theta(x_t, t) \right\|^2$$

$\mu_\theta(x_t, t)$ can be represented the same way as $\tilde{\mu}_t(x_t, x_0)$

$$\mu_\theta(x_t, t) = \frac{1}{\sqrt{\alpha_t}} \left(x_t - \frac{\beta_t}{\sqrt{1 - \bar{\alpha}_t}} \varepsilon_\theta(x_t, t) \right)$$

$$L_t = \frac{1}{2\sigma_t^2} \left\| \frac{1}{\sqrt{\alpha_t}} \left(x_t - \frac{\beta_t}{\sqrt{1 - \bar{\alpha}_t}} \varepsilon \right) - \frac{1}{\sqrt{\alpha_t}} \left(x_t - \frac{\beta_t}{\sqrt{1 - \bar{\alpha}_t}} \varepsilon_\theta(x_t, t) \right) \right\|^2$$



Reverse (Denoising) Process (14/15)

Simplification results in

$$L_t = \frac{\beta_t}{2\sigma_t^2\alpha_t(1-\bar{\alpha}_t)} \|\varepsilon - \varepsilon_\theta(x_t, t)\|^2$$

Scaling term

Actual noise

Predicted noise

DDPM authors found out that ignoring this term gives better training results

$$L_t = \|\varepsilon - \varepsilon_\theta(x_t, t)\|^2$$

$$= \|\varepsilon - \varepsilon_\theta(\sqrt{\bar{\alpha}_t}x_0 + \sqrt{1-\bar{\alpha}_t}\varepsilon, t)\|^2 \quad \text{Optimized during training}$$

Going back to the beginning

$$p_\theta(x_{t-1}|x_t) = \mathcal{N}(x_{t-1}; \mu_\theta(x_t, t), \Sigma_\theta(x_t, t))$$

$$\frac{1}{\sqrt{\alpha_t}} \left(x_t - \frac{\beta_t}{\sqrt{1-\bar{\alpha}_t}} \varepsilon_\theta(x_t, t) \right)$$

$\Sigma_\theta(x_t, t) = \sigma_t^2 \mathbf{I}$



Reverse (Denoising) Process (15/15)

Reparameterization trick: $\mathcal{N}(\mu, \sigma^2) = \mu + \sigma \odot z$

$$x_{t-1} = \frac{1}{\sqrt{\alpha_t}} \left(x_t - \frac{\beta_t}{\sqrt{1 - \bar{\alpha}_t}} \varepsilon_\theta(x_t, t) \right) + \sigma_t z \quad t > 1$$

One final term remains: $\log(p_\theta(x_0|x_1))$



$$\mathcal{N}(x_0; \mu_\theta(x_1, 1), \sigma_1^2 \mathbf{I})$$

The authors decided to sample this term noiselessly $z = 0$

$$x_{t-1} = \frac{1}{\sqrt{\alpha_t}} \left(x_t - \frac{\beta_t}{\sqrt{1 - \bar{\alpha}_t}} \varepsilon_\theta(x_t, t) \right) \quad t = 1$$



DDPM Training Process

Source image \mathbf{x}_0 is randomly sampled from the original data distribution $q(\mathbf{x}_0)$

t (time step) is sampled uniformly between 1 and T

Actual noise ϵ is sampled from a normal distribution

A neural network is trained via gradient descent to predict the noise ϵ_θ

Algorithm 1 Training

- 1: **repeat**
- 2: $\mathbf{x}_0 \sim q(\mathbf{x}_0)$
- 3: $t \sim \text{Uniform}(\{1, \dots, T\})$
- 4: $\epsilon \sim \mathcal{N}(\mathbf{0}, \mathbf{I})$
- 5: Take gradient descent step on
$$\nabla_{\theta} \|\epsilon - \epsilon_\theta(\sqrt{\bar{\alpha}_t}\mathbf{x}_0 + \sqrt{1 - \bar{\alpha}_t}\epsilon, t)\|^2$$
- 6: **until** converged

$$L_t = \|\epsilon - \epsilon_\theta(\sqrt{\bar{\alpha}_t}\mathbf{x}_0 + \sqrt{1 - \bar{\alpha}_t}\epsilon, t)\|^2$$

DDPM Sampling Process

Sampling starts from \mathbf{x}_T which is an *isotropic Gaussian distribution*

Neural network gradually denoises it until $t = 1$

A slightly less denoised image \mathbf{x}_{t-1} can be obtained using this equation

An image \mathbf{x}_0 is returned that looks similar to the original data distribution

Algorithm 2 Sampling

```
1:  $\mathbf{x}_T \sim \mathcal{N}(\mathbf{0}, \mathbf{I})$ 
2: for  $t = T, \dots, 1$  do
3:    $\mathbf{z} \sim \mathcal{N}(\mathbf{0}, \mathbf{I})$  if  $t > 1$ , else  $\mathbf{z} = \mathbf{0}$ 
4:    $\mathbf{x}_{t-1} = \frac{1}{\sqrt{\alpha_t}} \left( \mathbf{x}_t - \frac{1-\alpha_t}{\sqrt{1-\bar{\alpha}_t}} \boldsymbol{\epsilon}_\theta(\mathbf{x}_t, t) \right) + \sigma_t \mathbf{z}$ 
5: end for
6: return  $\mathbf{x}_0$ 
```

$$x_{t-1} = \frac{1}{\sqrt{\alpha_t}} \left(x_t - \frac{\beta_t}{\sqrt{1-\bar{\alpha}_t}} \varepsilon_\theta(x_t, t) \right) + \sigma_t z$$

$\beta_t = 1 - \alpha_t$




Neural Network (U-Net) Architecture

Position embeddings → To operate on a particular noise level

ResNet blocks → To use skip connections for merging the output of a previous layer with a layer ahead

Attention module → To allow a neural network to focus on a particular information at a time and ignore the rest

GroupNorm → To divide channels into groups and normalize features within each group. It works well for models with small batch sizes




A typical U-Net architecture

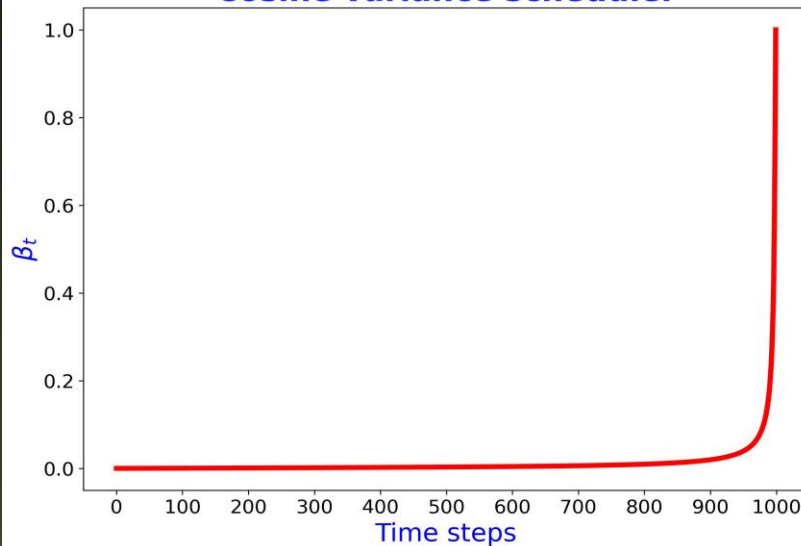


Different Variance Schedulers (1/2)


Linear variance scheduler



Cosine variance scheduler



Sigmoid variance scheduler



Linear




Cosine




Different Variance Schedulers (2/2)


Linear scheduler



Cosine scheduler



Sigmoid scheduler




$\bar{\alpha}_t$ plot for *linear*, *cosine* and *sigmoid* scheduler

These plots show how quickly or slowly the information in the source image is destroyed

We can easily observe that the information is destroyed much quicker in the case of *linear* than in cosine and sigmoid schedulers

Reverse Process Output Using Linear Scheduler




Reverse (denoising) process output

The final output should look like it came from the real data distribution



Random Sampling Using Linear Scheduler

Random noise



Randomly sampled images from noise



DDPM models will be able to generate actual images from noise only if trained well

- ◆ **The annotated diffusion model:** <https://huggingface.co/blog/annotated-diffusion>
- ◆ **What are diffusion models:** <https://lilianweng.github.io/posts/2021-07-11-diffusion-models/>
- ◆ **Denoising Diffusion Probabilistic Models:** <https://arxiv.org/pdf/2006.11239.pdf>
- ◆ **Improved Denoising Diffusion Probabilistic Models:** <https://arxiv.org/pdf/2102.09672.pdf>
- ◆ **U-Net Architecture:** <https://towardsdatascience.com/u-net-explained-understanding-its-image-segmentation-architecture-56e4842e313a>

THE END